

# CPSC 465/565 Theory of Distributed Systems



James Aspnes

2023-08-30

# Today's exciting topics

- ▶ Course overview
- ▶ Distributed computing
- ▶ Message passing model
- ▶ Safety, liveness, and fairness

# Course information

- ▶ Instructor: James Aspnes
- ▶ Teaching Fellows: John Lazarsfeld and Weijie Wang
- ▶ Course notes:  
<https://www.cs.yale.edu/homes/aspnes/classes/465/notes.pdf>
  - ▶ Also include lecture schedule, assignments, etc.
  - ▶ May change often.
- ▶ Coursework:
  - ▶ 5 assignments (100% of grade in 465, 85% in 565)
  - ▶ Presentation (565 only, 15% of grade)

Assignments will mostly involve proving things (it's a theory course).  
We will not be doing any programming or implementations.


# Distributed systems

“A distributed system is one where the failure of a computer you didn’t even know existed can render your own computer unusable.”


—Leslie Lamport

- ▶ Many processes
- ▶ Lots of **nondeterminism**, arising from
  - ▶ Message delays
  - ▶ Unpredictable scheduling
  - ▶ Failures

Do not confuse with parallel systems.



## Parallel systems are friendly



- ▶ Example: graphics cards
- ▶ Carefully synchronized processing units
- ▶ Predictable timing
- ▶ No failures

More processors = more power!

# Distributed systems are unfriendly



- ▶ Example: anything running over a network.
- ▶ Uncoordinated pile of machines
- ▶ Unpredictable timing
- ▶ Failures are normal and expected

More processes = more problems!

## Why theory?




Mathematical modeling + proofs are tools for dealing with nondeterminism.

- ▶ Exponentially many possible executions
- ▶ Not repeatable
- ▶ Bugs show up with low probability



# Reasoning about systems




- ▶ Model systems as mathematical objects
- ▶ Prove correctness
- ▶ Or prove impossibility

*$\forall$  executions; it works*





# Models




Distributed system =

- ▶ Collection of state machines
- ▶ Communications mechanism
  - ▶ Message passing: send packets across a network
  - ▶ Shared memory: read and write common address space
  - ▶ (or more exotic systems)

We will start by looking at message passing.



# Message passing




- ▶ Send messages across a network ( $\approx$  TCP/IP)
  - ▶ Network may or may not be complete
  - ▶ We will generally assume it is connected
- ▶ Nondeterminism:
  - ▶ Message delays
  - ▶ Message loss
  - ▶ Process failures:
    - ▶ Crash failures — process stops working
    - ▶ Byzantine failures — process turns evil!




# Shared memory



- ▶ Read and write shared objects
  - ▶ Often read/write registers (atomic registers)
  - ▶ Sometimes more powerful objects:
    - ▶ Test-and-set
    - ▶ Compare-and-swap
- ▶ Nondeterminism:
  - ▶ Asynchronous operations
  - ▶ Process failures (usually just crash failures)



## Typical distributed problems



- ▶ Agreement ←
- ▶ Replicated state machines
- ▶ Simulations, e.g. shared memory from message passing



# “The standard asynchronous model”

- ▶ Many historical candidates (c. 1970-1985 or so)
  - ▶ I/O automata
  - ▶ Temporal Logic of Actions
  - ▶ Communicating Sequential processes
  - ▶  $\pi$  calculus
  - ▶ pomsets
  - ▶ etc.
- ▶ Ultimately pretty much the same
- ▶ Compromise position:
  - ▶ I say I am using the standard model.
  - ▶ You assume I am using your favorite model.

We will use a model adapted from (Attiya and Welch, 2004).

# Message passing model



- ▶ Processes
  - ▶ Each process  $i$  has state in some state space  $Q$ .
- ▶ Buffers
  - ▶ Each pair of (connected) processes has buffer  $b_{ij}$ .
  - ▶ Buffer holds set of undelivered messages from  $M$ .
- ▶ Configuration = process states + buffer states
  - ▶ Formally, element of  $Q^n \times P(M)^{n \times n}$ .
- ▶ Transition function:
  - ▶  $\delta_i : Q \times P(M \times [n]) \rightarrow Q \times P(M \times [n])$ 
    - ▶ Old state  $\times$  messages delivered (with senders)
    - ▶  $\rightarrow$  new state  $\times$  messages sent (with recipients)



# Events



- ▶ Delivery event  $\text{del}(i, S)$  = deliver messages in  $S$  to  $i$ .
- ▶ Removes 0 or more messages from buffers.
- ▶ Applies transition function  $\delta$ :
  - ▶ Updates state of recipient
  - ▶ Adds 0 or more messages to recipient's outgoing buffers

## Execution



A **execution** is an alternating sequence of configurations  $C_t$  and events  $\alpha_t$ :

$$C_0 \alpha_1 C_1 \alpha_2 C_2 \alpha_3 C_3$$

- ▶ May be finite (ending in a configuration) or infinite.
- ▶ Constraints:

- ▶ Each  $\alpha_{t+1}$  is **enabled** in  $C_t$ :
  - ▶ Message passing:  $\text{del}(i, S)$  enabled if  $S \subseteq \cup_j b_{tj}$ .
  - ▶ Can only deliver messages that exist.
  - ▶ Other models will have different conditions.
- ▶  $C_{t+1}$  is the result of applying  $\alpha_{t+1}$  to  $C_t$

$$C \alpha \beta \gamma = C'$$

$$C_t \alpha_{t+1} = C_{t+1}$$



# Implicit assumptions in the message-passing model

- ▶ Only point-to-point messages
  - ▶ I can send messages to multiple processes at the same time.
  - ▶ But no guarantee they are delivered at the same time.
- ▶ Worse: “same time” only makes sense at one process:

$C_0\text{del}(p_1, S_1)C_1\text{del}(p_2, S_2)C_2$

- ▶ Impossible to represent simultaneous delivery events directly!
- ▶ (But not necessary.)


# Nondeterminism



A configuration  $C$  might have many enabled actions. Which to do?

- ▶ *All of them!*
- ▶ More specifically  $\forall$  possible executions, algorithm works.
- ▶ Universal quantifier anthropomorphized as the **adversary**.
- ▶ Adversary chooses which enabled action to do next.

## Safety and liveness



Want to prove an algorithm works no matter what the adversary does.

- ▶ Safety properties: nothing bad happens
- ▶ Liveness properties: something good happens eventually

Typically both are proved by induction over executions.

Liveness will require some constraints on adversary we'll discuss later.

## Example: flooding

Goal: Deliver a message to every process in an incomplete but connected network.

```
1 initially do
2   if pid = root then
3     received  $\leftarrow m$ 
4     send  $m$  to all neighbors
5   else
6     received  $\leftarrow \perp$ 
7 upon receiving  $m$  do
8   if received =  $\perp$  then
9     received  $\leftarrow m$ 
10    send  $m$  to all neighbors
```



# Proving correctness

Claim: Algorithm delivers  $m$  and only  $m$  to all processes.

Split into

1. In any reachable configuration,  $received_i \in \{\perp, m\}$ .
2. Eventually  $received_i \neq \perp$  forever.

where both claims hold for all processes  $i$ .

Note: (1) is a safety property, (2) is a liveness property.

future win  
↓

Eventually: 0 0 0 0 0 0

Forever: 0 0 0 0 0  
always win.

Ev P:  $\Diamond P$

For P:  $\Box P$

(2):  $\Diamond \Box received_i \neq \perp$

0 0 0 0 0 1 1 1 1 1  
always keep winning.

# Proving a safety property

Show  $P$  is an **invariant**:

- ▶  $P(C_0)$  holds.
- ▶ If  $P(C)$  holds and  $\alpha$  is enabled in  $C$ , then  $P(C\alpha)$  holds.

Since in any execution  $C_{t+1} = C_t\alpha$ , this gives

- ▶  $P(C_0)$
- ▶  $P(C_t) \Rightarrow P(C_{t+1})$

So induction gives  $P(C_t)$  for all  $t \in \mathbb{N}$ .


# Safety of flooding



Recall the claim:

In any reachable configuration,  $\text{received}_i \in \{\perp, m\}$ .

- ▶ Is it true in  $C_0$ ? Yes,  $\text{received}_i = m$  for initiator and  $\perp$  for everybody else.
- ▶ Does  $P(C) \Rightarrow P(C\alpha)$  always? Maybe...



## Using a stronger invariant

Let's try:

1. For all  $i$ ,  $\text{received}_i \in \{\perp, m\}$ .
2. For all  $ij$ ,  $b_{ij}$  contains only  $m$ .

Is this an invariant?

1. Is true in  $C_0$ . If  $\alpha$  changes  $\text{received}_i$  to  $m'$ , then  $m'$  was in some  $b_{ji}$  in a configuration satisfying (2). So  $m' = m$ .
2. Is true in  $C_0$ . If a new message  $m'$  is added to  $b_{ij}$ , it's because  $i$  received it from some  $b_{ki}$  in a configuration satisfying (2). So again  $m' = m$ .

(1+2) hold in init config.  
(1+2) preserved by  $\text{del}(i, m)$   
 $\alpha$




# Proving liveness

Eventually,  $\text{received}_i \neq \perp$  forever.


Since nothing in the code replaces  $\text{not-}\perp$  with  $\perp$ , it's enough to show that eventually  $\text{received}_i \neq \perp$  in some configuration.

But here is an execution in which this is not true:

$C_0 \text{del}(1, \emptyset) C_1 \text{del}(1, \emptyset) C_2 \text{del}(1, \emptyset) \dots$



# Fairness



To prove liveness, we must have **fairness**.

Define a **fair** execution by the rule

- Every message that is sent is eventually delivered. —

(Details may change in other models.)

Now insist only that fair executions satisfy liveness.

$$\begin{aligned} \forall \models, \text{works} \\ \forall \text{Fair} \models \text{works} \end{aligned}$$

# Liveness of flooding with fairness

Expand

- ▶ For all  $i$ , eventually received $_i \neq \perp$

to an induction hypothesis

- ▶ For all  $d$  and all  $i$  at distance  $d$  from the initiator, eventually  $i$  receives  $m$ .

Now do induction on  $d$ :

- ▶ Base case:  $d = 0$ . The initiator receives  $m$ .
- ▶ Induction step: Suppose  $i$  is at distance  $d + 1$ . Then  $i$  has at least one neighbor  $j$  at distance  $d$ . When  $j$  eventually receives  $m$  for the first time, it sends  $m$  to  $i \Rightarrow$  eventually  $i$  receives  $m$ .

# How much does this cost?

Some complexity measures are straightforward:

- ▶ Local computation: How much work did each process do?
  - ▶ Don't care. Pretend it's free.
- ▶ Message complexity: How many messages were sent?
  - ▶ Flooding sends  $2|E|$  messages.
- ▶ Bit complexity: What was the total size of those messages?
  - ▶ Flooding sends  $2|E||m|$  bits.

Some are not:

- ▶ Time complexity: How much time until algorithm finishes?
  - ▶ Maybe a long time with unbounded message delays!
  - ▶ We'll revisit this next time.

