

# CPSC 465/565 Theory of Distributed Systems

James Aspnes

2023-09-01

# Today's exciting topics

- ▶ Message passing complexity measures
- ▶ Broadcast and convergecast

# Message passing model



- ▶ Buffer  $b_{ij}$  holds messages from  $p_i$  to  $p_j$
- ▶ Execution:  $C_0\alpha_1C_1\alpha_2C_2\dots$ 
  - ▶  $\alpha_t$  = delivery event at step  $t$
- ▶ Nondeterminism: adversary chooses next event from enabled events
- ▶ Fairness: Every message sent is eventually received



What about complexity?


# Complexity measures for message-passing

Last time:

- ▶ Local computation: free!
  - ▶ Justification: usually fast relative to message delivery times
- ▶ Message complexity = number of messages sent
- ▶ Bit complexity = total size of all messages

This time:

- ▶ Time complexity = time to finish
  - ▶ But: even in fair executions, no upper bound on delivery time



## Time complexity without time bounds




- ▶ Define maximum message delay = 1 time unit.
- ▶ Define gap between two local computation events = 0 time units.
- ▶ Assign each event the max time consistent with these assumptions.

This gives a time measure that is agnostic about actual time delays.

(Similar to assumption behind  $O()$  notation.)

## An asynchronous execution



Assign highest possible time to each event consistent with:

- ▶  $T(\text{first step}) = 0$ .
- ▶  $T(\text{receive}) \leq T(\text{send}) + 1$ .
- ▶ Consecutive local computation events have same time.


## Example: flooding with parent pointers

```
1 initially do
2   if pid = root then
3     parent  $\leftarrow$  root
4     send  $m$  to all neighbors
5   else
6     parent  $\leftarrow \perp$ 
7 upon receiving  $m$  from  $j$  do
8   if parent =  $\perp$  then
9     parent  $\leftarrow p_j$ 
10    send  $m$  to all neighbors
11
```



Builds a spanning tree with broadcast as side-effect.

## Correctness



Claim: eventually parent pointers form rooted spanning tree.

► Safety:


1.  $\text{parent}_i = \perp$  or  $\text{parent}_i$  is node on path to root
2. If  $m \in b_{ij}$ , then  $\text{parent}_i \neq \perp$ .

► Liveness: Eventually  $\text{parent}_i \neq \perp$ .



# Safety proof

1.  $\text{parent}_i = \perp$  or  $\text{parent}_i$  is node on path to root
  2. If  $m \in b_{ij}$ , then  $\text{parent}_i \neq \perp$ .
- ▶ Check initial configuration:
    1.  $\text{parent}_i = \perp$  unless  $i = \text{root}$ ; count  $\text{parent}_{\text{root}} = \text{root}$  as OK.
    2. If  $m \in b_{ij}$ , then  $i = \text{root}$  and  $\text{parent}_{\text{root}} \neq \perp$ .
  - ▶ Check each claim against steps that could make it false:
    1. If  $i$  sets  $\text{parent}_i$  to  $j$ :
      - ▶  $i$  receives  $m$  from  $j$
      - ▶  $m \in b_{ij} \Rightarrow \text{parent}_j \neq \perp \Rightarrow \text{parent}_j$  is on path to root
    2. If  $i$  sends  $m$ :
      - ▶ Same code sets  $\text{parent}_i \neq \perp$




# Liveness proof

Eventually  $\text{parent}_i \neq \perp$ .

Same induction on distance as simple flooding.

But now we will include time.



- ▶ Claim: For all  $i$ ,  $\text{parent}_i \neq \perp$  by time  $d(i, \text{root})$ .
- ▶ Proof: By induction on  $d(i, \text{root})$ :
  - ▶ Base:  $d(i, \text{root}) = 0 \Rightarrow i = \text{root}$ , so  $\text{parent}_i = i$  at time 0.
  - ▶ Induction step: Suppose  $d(i, \text{root}) = d + 1$ 
    - ▶ Then there is some neighbor  $j$  that sets  $\text{parent}_j \neq \perp$  by time  $d$ .
    - ▶ When  $j$  does this, it sends  $m$  to  $i$ .
    - ▶ At most one time unit later ( $\leq d + 1$ ),  $i$  receives  $m$ .
    - ▶ When  $i$  receives  $m$ , sets  $\text{parent}_i \neq \perp$  if not already set.



## Time complexity

$D = \text{diameter}$   
 $= \max \text{ dist}$   
between 2 nodes

Count until last message delivered.

- ▶ Most distant node is at  $\max_i d(i, \text{root}) \leq \text{diam}(G)$ .
- ▶ Every node sets  $\text{parent}_i \neq \perp$  by time  $D$ .
- ▶ Last message is sent by time  $D$ .
- ▶ Last message is received by time  $D + 1$ .

So time complexity is  $D + 1$ .

# Detecting termination


Add acknowledgments:

```
1 upon receiving  $m$  from  $j$  do  
2   if  $\text{parent} = \perp$  then  
3      $\text{parent} \leftarrow j$   
4     send  $m$  to all neighbors  
5   else  
6     send nack to  $j$   
7 upon receiving ack from  $j$  do  
8    $\text{children} \leftarrow \text{children} \cup \{j\}$   
9 upon receiving nack from  $j$  do  
10   $\text{nonChildren} \leftarrow \text{nonChildren} \cup \{j\}$   
11 as soon as  $\text{children} \cup \text{nonChildren} = \text{neighbors}$  do  
12  send ack to parent
```



Done when root gets ack or nack from all neighbors.

# Complexity of the terminating version



- ▶ Message complexity


- ▶ Every edge gets  $m$  in each direction
- ▶ Plus ack or nack in response
- ▶ Total is  $4|E|$  messages.

$2|E|$   
receiving  
msgs

- ▶ Time complexity

- ▶ Outgoing messages reach everybody in time  $D$ .
- ▶ All nacks are sent by time  $D + 1$ .
- ▶ All nacks are received by time  $D + 2$ .
- ▶ What about acks?

$+2|E|$   
ack to nacks




# Timing of acks

Look at position of each node in the tree:


- ▶ If I am a leaf:
  - ▶ I send ack once I collect all nacks
  - ▶ This happens by time  $D + 2$
- ▶ If I am a parent:
  - ▶ I have to wait for all my kids.

Induction hypothesis:

- ▶ Node at height  $h$  sends ack at time  $\leq D + 2 + h$ .
- ▶  $\Rightarrow$  finish by time  $D + 2 + (\text{height of tree})$ .



## How tall can the tree get?

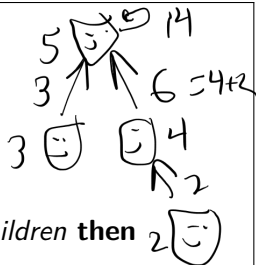


- ▶ Max distance from root is 1, so all nodes recruited by time 1.
- ▶ But maybe some nodes are recruited faster.
- ▶ Worst case: height of tree =  $n - 1$ .
- ▶ Gives  $D + 2 + (n - 1)$  worst-case time for broadcast with acks.

# Convergecast

Sometimes it's convenient to collect data from an existing tree:

```
1 initially do
2   if I am a leaf then
3     send input to parent
4 upon receiving  $v_c$  from  $c$  do
5   append  $(c, v_c)$  to buffer
6   if buffer contains messages from all my children then
7      $v \leftarrow f(\text{buffer}, \text{input})$ 
8     if pid = root then
9       return  $v$ 
10    else
11      send  $v$  to parent
```




**Broadcast** to build the tree and **convergecast** to collect data.



# Performance of convergecast

- ▶ Message complexity =  $n - 1$
- ▶ Time complexity:
  - ▶ Leaves send at time 0
  - ▶ Parent sends by max child time +1
  - ▶  $\Rightarrow$  each  $i$  sends by  $\text{height}(i)$
  - ▶  $O(n)$  time in worst case.

# Synchronous model




Assumption: every message delay = 1 exactly.

Formal version:

- ▶ Sequence of rounds  $r = 0, 1, 2, \dots$
- ▶ Every node gets one delivery event per round.
- ▶ Every message sent in  $r$  is received in  $r + 1$ .

## Synchronous broadcast gives shortest-path tree



Proof: By induction,  $i$  receives  $m$  exactly at time  $d(\text{root}, i)$ .

~~end{center}~~


# What if we don't have synchrony?

Attempt 1: Use Dijkstra's algorithm.

```
1 initially do
2   if pid = initiator then
3     dist  $\leftarrow$  0
4     send dist to all neighbors
5   else
6     dist  $\leftarrow \infty$ 
7 upon receiving  $d$  from  $p$  do
8   if  $d + 1 < \text{dist}$  then
9     dist  $\leftarrow d + 1$ 
10    parent  $\leftarrow p$ 
11    send dist to all neighbors
```



## Dijkstra in action



Claim: Every process converges to correct distance.

- Safety:  $\text{dist}_i \geq d(i, \text{root})$  and for any  $d \in b_{ij}$ ,  $d \geq d(i, \text{root})$ .
- Liveness:  $\text{dist}_i = d(i, \text{root})$  no later than time  $d(i, \text{root})$ .

Ind on  $d$ : Base: true for  $d=0$


Ind step: closer nbr at dist  $d$  sets by time  $d$   
 $\Rightarrow$  msg  $\Rightarrow$  I set to  $d+1$  by time  $d+1$ .

## Worst-case complexity

↙ diameter

Liveness proof gives  $O(D)$  time.

But message complexity may be bad:



- Careful schedule yields  $\Theta(n^2)$  updates.
- $\Rightarrow \Theta(n^3)$  messages.
- With synchrony, only  $2|E| = O(n^2)$  messages.

## Fix by simulating synchrony


( $\propto$  synchronizer)

```
1 initially do
2   if pid = root then
3     dist  $\leftarrow$  0
4     send [= 0] to all neighbors
5   else
6     send [ $\geq$  0] to all neighbors
7 as soon as I received [ $\geq d$ ] from all neighbors do
8   send [ $\geq d + 1$ ] to all neighbors
9 as soon as I received [= d] from at least one neighbor and
  [ $\geq d$ ] from the others do
10  dist  $\leftarrow d + 1$ 
11  send [= d + 1] to all neighbors
12  halt
```

$(= \downarrow) \rightarrow$  I am exactly

$(\geq \downarrow) \rightarrow$  I am at least  $d$

# Correctness



## ► Safety:

- ✓ ► If  $i$  sends  $[\geq d]$ ,  $d(i, \text{root}) \geq d$ .
- If  $i$  sends  $[= d]$ ,  $d(i, \text{root}) = d$ .

+ similar constraints on msgs.

## ► Liveness:

- If  $d(i, \text{root}) = d$ , then
  - For each  $d' < d$ ,  $i$  sends  $[\geq d']$  by time  $d'$ .
  - $i$  sends  $[= d]$  by time  $d$ .

Easy induction on  $d$ .  
(rhymes w/ lazy)






# Cost

- ▶ From safety property:
  - ▶ Distance- $d$  node sends only  $[\geq 0], [\geq 1], \dots, [\geq d-1], [=d]$ .
  - ▶ This gives at most  $D+1$  messages per direction per edge.
  - ▶  $\Rightarrow$  Message complexity =  $O(D|E|)$ .
- ▶ Liveness property:
  - ▶ Most distant node sends  $[=d]$  by time  $d$ .
  - ▶  $\Rightarrow$  Time complexity =  $O(D)$ .

For example, get  $O(n^2)$  messages and  $O(1)$  time on  $K_n$ .



Next time

Leader election!