


CPSC 465/565 Theory of Distributed Systems

James Aspnes

2023-09-18



Today's exciting topics

Synchronous agreement with failures, including:


- ▶ Coordinated attack!
- ▶ Upper and lower bounds on time!
- ▶ Byzantine agreement!

Two generals problem (Gray 1978, Akkoyunlu et al. 1975)




- ▶ Two generals separated by a dangerous enemy camp.
- ▶ Attack or retreat?
 - ▶ Both attack: Victory!
 - ▶ Both retreat: Maybe victory tomorrow.
 - ▶ One attacks, one retreats: Shameful defeat!
- ▶ Model: message-passing with omission failures
- ▶ Question: Can we reach agreement always?

Formal version



- ▶ **Coordinated attack problem** (n -general version)
- ▶ Synchronous message passing with omission failures.
 - ▶ Any message can be lost.
 - ▶ Sender doesn't know about lost messages.
- ▶ Requirements:
 - ▶ **Agreement:** All processes output same 0-1 value.
 - ▶ **Termination:** Protocol finishes in bounded rounds.
 - ▶ **Validity:**
 - ▶ If all processes have the same input
 - ▶ and no messages are lost,
 - ▶ all processes output the common input.




Impossibility of coordinated attack

$\approx \sim_p \Xi'$

Claim: No protocol survives unlimited message losses.


iff $\Xi \models (p \models \Xi) \wedge (q \models \Xi)$



Proof:

- ▶ Fix a supposedly correct protocol with processes p and q .
- ▶ Consider execution Ξ_0 with all 0 inputs and no lost messages.
 - ▶ Validity \Rightarrow both processes decide 0.
 - ▶ Termination \Rightarrow finite sequence of messages.
- ▶ What happens if we delete last message from p to q ?
 - ▶ New execution $\Xi_0^{-1} \sim_p \Xi_0$.
 - ▶ So p still decides 0.
 - ▶ Agreement $\Rightarrow q$ still decides 0.
- ▶ Now delete last message from q to p to get Ξ_0^{-2} .
 - ▶ Now $\Xi_0^{-2} \sim_q \Xi_0^{-1}$.
 - ▶ So both p and q decide 0 in Ξ_0^{-2} .

Impossibility of coordinated attack




Remove messages one at a time until all messages are lost.

- ▶ Get $\Xi_0 \sim_p \Xi_0^{-1} \sim_q \Xi_0^{-2} \sim_p \Xi_0^{-3} \sim_q \dots \Xi_0^{-m}$, where
 - ▶ Each Ξ_0^{-k} has both processes output 0, and
 - ▶ Ξ_0^{-m} delivers no messages.

Now do the same thing starting with Ξ_1 with 1 inputs:

- ▶ Get $\Xi_1 \sim_p \Xi_1^{-1} \sim_q \Xi_1^{-2} \sim_p \Xi_1^{-3} \sim_q \dots \Xi_1^{-m'}$, where
 - ▶ Each Ξ_1^{-k} has both processes output 0, and
 - ▶ $\Xi_1^{-m'}$ delivers no messages.

Impossibility of coordinated attack



1
6

Finally, construct Ξ_{01} where

- ▶ p has input 0
- ▶ q has input 1
- ▶ All messages are lost, so $\Xi_0^{-m} \sim_p \Xi_{01} \sim_q \Xi_1^{-m'}$.



Since p decides 0 in Ξ_0^{-m} , p decides 0 in Ξ_{01} .


Since q decides 1 in $\Xi_1^{-m'}$, q decides 1 in Ξ_{01} .

So Ξ_{01} violates agreement.

(For $n > 2$, same argument, just lose more messages.)

Synchronous model with crash failures

Crash failure: Faulty process dies!



- ▶ Some subset of messages in crash round are delivered.
- ▶ No messages in subsequent rounds are delivered.
- ▶ $f \leq t$ processes fail in an execution.
 - ▶ f = actual failures
 - ▶ t = maximum failures (“tolerance”)

This is a stronger model than omission failures.

Synchronous agreement with crash failures

Almost the same requirements as coordinated attack:


- ▶ **Agreement:** All non-faulty processes output same value.
- ▶ **Termination:** Protocol finishes in bounded rounds.
- ▶ **Validity:** All inputs equal \Rightarrow all non-faulty processes output input.
 - ▶ Equivalent version: common output is somebody's input.
 - ▶ We'll require this even if there are failures.

With crash failures this turns out to be possible.

Dolev-Strong 1983

Synchronous agreement with up to t crash failures in $t + 1$ time.

```
1  $S_i \leftarrow \{\langle i, \text{input}_i \rangle\}$  //  $S$  stores known id-input pairs
2 for  $r \leftarrow 1$  to  $t + 1$  do
3   Send  $S_i$  to all processes (including myself)
4   Receive messages  $S_k$ 
5    $S_i \leftarrow \bigcup S_k$ 
6 return  $f(S)$ 
```



This works for any reasonable f (smallest input, smallest id, ...).

Note f is same for all processes.

Easy properties:

- ▶ Termination: Loop only runs for $t + 1$ rounds.
- ▶ Validity: f picks a value that started off as somebody's input.


Dolev-Strong: agreement

Agreement requires a bit more of an argument.


```
1  $S_i \leftarrow \{\langle i, \text{input}_i \rangle\}$  //  $S$  stores known id-input pairs
2 for  $r \leftarrow 1$  to  $t + 1$  do
3   Send  $S_i$  to all processes (including myself)
4   Receive messages  $S_k$ 
5    $S_i \leftarrow \bigcup S_k$ 
6 return  $f(S)$ 
```



- ▶ By Pigeonhole Principle, \exists a round s with no new failures.
- ▶ Let $S_i^r =$ value of S_i after r rounds.
- ▶ Claim: $S_i^s = S_j^s$ for all non-faulty i and j .
 - ▶ Proof: i and j receive the same messages in round s .
- ▶ Claim: $S_i^r = S_j^r$ for all non-faulty i, j and all $r \geq s$.
 - ▶ Proof: Induction on r : $S_i^{r+1} = \bigcup S_k^r = S_i^r = S_j^r$.
- ▶ $\Rightarrow f(S_i^{t+1}) = f(S_j^{t+1})$.



Are $t + 1$ rounds necessary?




For the Dolev-Strong algorithm, yes:

- ▶ Suppose $f(S) = \max$.
- ▶ Start 1 with input 1, all others with 0.
- ▶ In each round r , crash process r and deliver S_r to $r + 1$ only.

One process sees $\{0, 1\}$ and decides 1, while others decide 0.

Are $t + 1$ rounds necessary for any algorithm?

Yes! The lower bound is due to (Dolev and Strong, 1985).




$$\begin{matrix} 0 \\ 0 \\ 0 \end{matrix} \rightsquigarrow \begin{matrix} 1 \\ 1 \\ 1 \end{matrix}$$

Handwritten notation illustrating a transformation or sequence of states, possibly representing a sequence of inputs or outputs.

Strategy:


- ▶ Start with all-0 input execution.
- ▶ Build a chain of intermediate executions to change a 0 to a 1.
- ▶ Exploit indistinguishability to make each step in the chain preserve outputs.
- ▶ Arrive at all-1 input execution deciding 0.

How to hide a change



- ▶ Want to change input at process p .
- ▶ If we crash p , nobody else will know!
- ▶ But maybe they saw we crashed p .
 - ▶ Solution: remove p 's outgoing messages one at a time.
 - ▶ For each removal, crash the receiver first.
- ▶ But then we see the receiver crashed!
 - ▶ Solution: Slaughter witnesses recursively.

Main lemma




Claim: Let Ξ_0 be a t -round execution in which:

- ▶ At most one process crashes in rounds $1..r-1$.
- ▶ No process crashes in rounds $r..t$.

Then for any non-faulty q in Ξ_0 , there is a sequence $\Xi_0, \Xi_1, \Xi_2, \dots, \Xi_m$ such that:

- ▶ $\Xi_i \sim_p \Xi_{i+1}$ for some p that doesn't crash in either execution.
- ▶ Ξ_m has the same message pattern as Ξ_0 , except
- ▶ q crashes fully (sends no messages) in round r in Ξ_m .

Proof of main lemma



By induction on $t - r$:

1. If $r = t$, remove q 's outgoing messages one at a time.
 - Only the former recipient notices each step.
2. If $r < t$, pick a non-faulty recipient s .
 - 2.1 Use the lemma recursively to crash s .
 - 2.2 Remove the message to s .
 - 2.3 Use the lemma in reverse to uncrash s .
 - 2.4 Repeat until q has no outgoing messages in rounds r or higher.

Rest of proof



Given a t -round protocol that supposedly solves consensus:


1. Use main lemma to crash a process at start of protocol.
2. Change its input from 0 to 1.
3. Repeat until all inputs changed.

Somewhere in this exponentially long sequence of executions, we violate agreement or violate validity.

\Rightarrow Need at least $t + 1$ rounds.

(This is tight because Dolev-Strong algorithm uses $t + 1$ rounds.)

Byzantine failures



Worse than mere death: processes turn evil!

Byzantine agreement (Pease, Shostak, Lamport 1980).

- ▶ A Byzantine process can send any message it likes.
 - ▶ Allied to the adversary
 - ▶ Not bound by the protocol.
 - ▶ Seeks only to destroy.
- ▶ Constraints
 - ▶ Can't impersonate other processes.
 - ▶ Model is still synchronous.
- ▶ Name is offensive to Byzantine Empire (ended 1453).
 - ▶ May also annoy some Byzantine Orthodox Christians.
 - ▶ We are kind of stuck with it.

Where is your agreement, termination, and validity now?


Requirements for Byzantine agreement

Revised to take into account evil nodes:

- ▶ Agreement: All non-faulty nodes output same value.
- ▶ Termination: Finish in bounded number of rounds.
- ▶ Validity: If all *non-faulty* nodes have same input, they all output this input.

Validity in particular means Byzantines can't hijack the protocol.

Impossibility results





Dolev-Strong implies $t + 1$ rounds needed for t Byzantine failures.

Pease et al. show $n \geq 3t + 1$ also required:

- ▶ Intuition: With 1 evil node out of 3 total nodes, evil node can play good nodes against each other.
- ▶ Generalizes to $t = n/3$ since each of these nodes can simulate $n/3$ nodes.
- ▶ Need more than a story to get a real proof.

The diagrams illustrate the following processes:

- Mitosis:** A rectangular cell containing four red dots (chromosomes) divides into two identical daughter cells, each containing two red dots.
- Meiosis:** A rectangular cell containing four blue dots (chromosomes) divides into two daughter cells. Each daughter cell then divides again, resulting in four daughter cells, each containing two blue dots.
- Binary Fission:** A triangular cell containing two dots (one red, one blue) divides into two daughter cells, each containing one dot.



(Fischer, Lynch, Merritt 1986)

- ▶ Adversary simulates six non-faulty processes in a ring.
- ▶ Each copy X_i has input i .
- ▶ A_1 and B_1 think they are in an execution with evil C .
 - ▶ Validity means A_1 and B_1 both decide 1.
 - ▶ In general all X_i decide i .
- ▶ Now run evil B against good A_1 and C_0 .
 - ▶ B acts like B_1 with A_1 and B_0 with C_0 .
 - ▶ As in 6-process execution, A_1 decides 1, C_0 decides 0.
 - ▶ \Rightarrow No agreement!

► \Rightarrow No agreement!

Next time

- ▶ Protocols for synchronous Byzantine agreement ($n \geq 3t + 1$).
- ▶ Impossibility of asynchronous agreement with one crash failure.

